

# Energy Aggregation using Product of Experts\*

Megha Gupta [meghag@iiitd.ac.in](mailto:meghag@iiitd.ac.in)

Haimonti Dutta [haimonti@buffalo.edu](mailto:haimonti@buffalo.edu)

Amarjeet Singh [amarjeet@iiitd.ac.in](mailto:amarjeet@iiitd.ac.in)

Ullas Nambiar [Ullas.Nambiar@emc.com](mailto:Ullas.Nambiar@emc.com)

## Abstract

## 1 Introduction

Smart meters consisting of real time sensors, power outage notifications and power quality monitoring are widely used today. These meters provide a host of benefits like energy efficiency and savings, improved retail competition, better demand response actions, improved tariffs, lower bills due to better customer feedback, accurate billing, less environmental pollution, etc. They generate huge amount of data which helps in giving meaningful insights through analytics. They can measure site specific information and also help agencies to set different electricity prices for consumption based on the time of the day, seasons, holidays, etc. As a result, a feedback is sent to the customers by the utilities that can help consumers better manage their resources. A research [McKerracher and Torriti, 2013] shows that by providing real time feedback, consumers can reduce the consumption by 3-5%. Also, for some country providing real time feedback may not be a cost effective plan but it can help in retaining customers and peak-load shift.

In recent years, machine learning has been applied to the problem of energy consumption and demand forecasting analysis. The role of the machine learning algorithm is to study the sensor data and provide alerts and warnings when anomalous behaviour occurs or to inform (and remind) customers when certain activities were performed, which rooms they occupied, and what appliances they used most frequently during that period. This information can be transmitted to customers in timely fashion via phone, email or the Internet. This paper [Chicco *et al.*, 2006] does a comparison of several clustering techniques and finds out that the hierarchical clustering and modified follow-the-leader perform best among the rest K-Means, fuzzy K-Means to group customers with similar electrical behaviour [Martinez Alvarez *et al.*, 2011]. Another paper [Wijaya *et al.*, 2013] uses classifiers like random forest, J48, logistic and naive bayes to identify customers with similar electricity consumption profiles. Sensor data collected from smart homes are used to reveal activity patterns

of the residents, which can then be correlated with the total energy consumption. This enables utility companies and their customers to associate activities with energy usage and costs, devise intelligent systems to control home environments improving energy efficiency and reducing costs. Typically, sequences of usage patterns that appear frequently at different time scales (daily, weekly, monthly, yearly) and across different homes are studied and outlier detection algorithms are designed to enable customers to be notified that they are consuming unusually large amount(s) of energy during some specific period. Related problems involve study of trends of electricity consumption (steadily increasing, decreasing, cyclic, seasonal) and sudden anomalous behaviour (sudden peaks or drops on consumption) for individual homes and across the community.

In this paper, we build machine learning models using products of HMM and apply them to the energy aggregation problem. Two different proof of concepts are presented – one on the REDD data set and another on real data collected at the faculty housing in India. There are many reasons why the product model constructed from many HMMs is appropriate. First, this model is ideal for data which is caused by multiple underlying influences. Second, HMMs alone are not efficient at capturing long range structure in time series [Taylor and Hinton, 2009] – in contrast to product of hidden markov models (PoHMM) [Brown, 2001] allow each model to remember a different piece of information about the past.

**Organization:** This paper is organized as follows: Section 2 examines related work on data analytics on aggregated data of smart meters; Section 3 provides a review of products of Hidden Markov Models (HMMs) and how they relate to our application. The two proofs of concepts are introduced in Section 4 to illustrate the effectiveness of the use of product of HMMs in the energy aggregation problem. Finally, Section 6 concludes the work.

## 2 Related Work

### 2.1 Energy Aggregation

Data aggregation also known as data fusion is defined as set of methods for combining the data that comes from many sensor nodes into a set of meaningful information [Heinzelman *et al.*, 2000]. To study the process of energy aggregation, data is collected from multiple smart meters. This collected data is

---

\*These match the formatting instructions of IJCAI-07. The support of IJCAI, Inc. is acknowledged.

very large which in turn makes the analytics on top of it very difficult. Also, this detailed energy consumption data leads to privacy breach and other risks related to it. To address this problem, there has been work done [Wijaya *et al.*, 2013] to reduce the smart meter data numerosity by converting it into symbolic representation and then allowing algorithms on top of it. They have applied the symbolic representation tasks for customer segmentation and load forecasting.

## 2.2 Energy Disaggregation

The process in which the whole building energy (aggregated) signal is separated into appliance level energy (disaggregated) for a variety of reasons like residential energy reductions, program evaluation, targeted marketing, etc. Several studies have been done in this regard, one of the unsupervised desegregation method [Kim *et al.*, 2011] that outperforms other unsupervised disaggregation methods is conditional factorial hidden semi-Markov model. This model when integrated with other features, accurately represents the individual appliance energy consumption. Another research [Kolter and Jaakkola, 2012] that exploits the additive structure of the FHMM to develop approximate inference procedure in energy disaggregation domain that outperforms the rest.

## 2.3 Load Forecasting

Load forecasting is extremely important for operation and planning of an electric utility company. It is usually divided into three categories, short-term forecasts [Bakirtzis *et al.*, 1996] [Chen *et al.*, 2001] [Fan and McDonald, 1994] [Chow and Leung, 1996], medium-term forecasts [Falvo *et al.*, 2006] [Samuel *et al.*, 2014] [Bassi and Olivares, 2006] and long-term forecasts [Daneshi *et al.*, 2008] [Zhang and Ye, 2011]. Most of the load forecasting algorithms use statistical and artificial intelligence algorithms.

## 2.4 Customer Segmentation

The identification of consumer profiles that show similar behaviour in energy consumption. This analysis is useful in various ways, like demand response system, intelligent distribution channel. The author [Wijaya *et al.*, 2014] segments the customers based on contextual dimensions like location, seasons, weather patterns, holidays, etc which help with various higher level applications like usage-specific tariff structure, theft detection, etc. In [Albert and Rajagopal, 2013], author proposes to infer occupancy states from consumption time series data by using HMM framework. They investigate the effectiveness of HMM and model based cluster analysis in producing meaningful features of the classification. This work suggests the dynamics of time series as captured by HMM analysis can be valuable.

## 2.5 Prediction with expert advice

## 3 Review of PoHMM

Distributed networks can be modeled using interacting automata. Benveniste [Fabre *et al.*, 2000] defines automaton as a quadruple,  $\hat{A} = (X, X_0, A, T)$  where  $X$  is a finite state of sets,  $X_0$  is the subset of initial states,  $A$  is a finite set of messages,  $T$  is a set of transitions of the form  $t = \{x_-, a, x\}$

where  $x_-$  is the previous state,  $a$  is the message label on which the state transitions to the next state  $x$ . The figure 1 below explains the automata with an example.

For automaton  $R$ ,  $X_R = \{2; R_1, R_2\}$ ,  $X_{0R} = \{R_1\}$ ,  $A_R = \{a, b\}$ ,  $T_R = \{R_1, a, R_1; R_1, b, R_2; R_2, a, R_2; R_2, b, R_1\}$   
For automaton  $S$ ,  $X_S = \{3; S_1, S_2, S_3\}$ ,  $X_{0S} = \{S_1\}$ ,  $A_S = \{a, b\}$ ,  $T_S = \{S_1, a, S_1; S_1, b, S_2; S_2, a, S_2; S_2, b, S_3; S_3, a, S_3; S_3, b, S_1\}$

The product of two automata  $\hat{A} = R \times S$  is defined as follows:  
 $X = X_R \times X_S$   
 $X_0 = X_{0R} \times X_{0S}$   
 $A = A_R \cup A_S$

Benveniste uses a notion of stuttering transition which helps to distinguish between local and global time by inserting dummy transitions between two transitions of a local automaton attached to a node. This stuttering transition waits for others to progress.

A	$R_1$	$R_2$
$R_1$	0.6	0.4
$R_2$	0.3	0.7

Table 1: Transition probability, A

B	a	b
$R_1$	0.2	0.8
$R_2$	0.5	0.5

Table 2: Observed probability, B

	$R_1$	$R_2$
$\pi$	0.4	0.6

Table 3: Initial state probability,  $\pi$

Talking in terms of HMM, requires us to equip products of automata with probabilities. Benveniste defines HMM as a triple  $(\hat{A}, \mu, \pi)$  where  $\hat{A} = (X, X_0, A, T)$  is an automaton,  $\mu$  is the initial state probability,  $\pi$  is factored as state transition probability  $\pi_x$  and message transition probability  $\pi_A$ . He uses a random arbiter  $\alpha$ , with values first, second, third to choose automaton to initiate transition. If  $\alpha = \text{first}$  then first automaton chooses any transition having a private message whereas second automaton performs a stuttering transition, and vice versa for  $\alpha = \text{second}$ . If  $\alpha = \text{both}$ , then both automata agree on some shared message and move accordingly.

Using the traditional HMM notation of the parameters  $\lambda = \{A, B, \pi\}$  where  $A$  is the transition probability,  $B$  is the observed probability,  $\pi$  is the initial state probability. For automata  $R$ , we have the values of  $A, B, \pi$  as shown in table 1, 2, 3 respectively.

### 3.1 Product of HMM

PoHMM is a combination of directed and undirected graphical models. The hidden states are connected with directed links where as the connection with visible states is




Figure 1: Automata R and S

through undirected links. This causes different conditional independence relationships among the variables in graphical model. Product of HMM is a way of combining HMM's to form distributed state time series model. It is defined by multiplying together the densities of its,  $k$  experts and renormalizing them. The figure 2 is a product of two HMMs shown in 1. For  $P = R \times S$ , the quadruple becomes  $X = \{6; R_1S_1, R_1S_2, R_1S_3, R_2S_1, R_2S_2, R_2S_3\}$   
 $X_0 = \{R_1S_1\}$   
 $A = \{a, b\}$

The rules for synchronised product construction are :

1.  $\langle p, q \rangle \xrightarrow{-a} \langle p', q' \rangle$  if  $a \in A_R \cap A_S$  and  $p \xrightarrow{-a} p'$  and  $q \xrightarrow{-a} q'$
2.  $\langle p, q \rangle \xrightarrow{-a} \langle p', q' \rangle$  if  $a \in A_R, a \notin A_S$  and  $p \xrightarrow{-a} p'$
3.  $\langle p, q \rangle \xrightarrow{-a} \langle p, q' \rangle$  if  $a \notin A_R, a \in A_S$  and  $q \xrightarrow{-a} q'$

### 3.2 Inference in PoHMM

The main feature of PoE is its undirected graphical modelling with no direct connection among the latent variables as they only interact indirectly via observed variables. The hidden variables all the experts are rendered independent when conditioned on visible variables. So, if the inference in each of the constituent model is tractable then the inference in the product is also tractable. To generate a data point in this model, all the experts in PoE generate an observation and if they all generated the same point then it is accepted else they again generate an observation until all the experts agree to it. Therefore all the experts have some influence over the generated data. So, the inference determines the the probability that all the experts would have taken in order to generate the given observation.

### 3.3 Training product of experts by minimising contrastive divergence

PoE is a method of combining densities of many latent variable models. It is defined by the following formula:

To fit the model to the data, we need to maximize the likelihood of the dataset or minimise the Kullback-Liebler divergence between the real data and the fantasy data. The contrastive divergence algorithm for training the PoHMM has the following steps:

1. Calculate each model's gradient on a data point using forward backward algorithm.
2. For each model take a sample from the posterior distribution of paths through state space.

3. At each time step, multiply together the distributions and renormalize to get the reconstruction distribution at each step.
4. Draw a sample from the reconstruction distribution at each time step to get a reconstructed sequence. Compute each model's gradient on the new sequence.
5. Update the parameters

## 4 Proof of Concepts

### 4.1 REDD House 2

#### 4.2 Aim

To represent streams of energy consumption data from  $n^1$  appliances by product of  $k$  HMMs.

#### 4.3 Method

- **Data** The Reference Energy Disaggregation Data Set (REDD) is used in empirical analysis. The data contains power consumption from real homes, for the whole house as well as for each individual circuit in the house (labeled by the main type of appliance on that circuit). It is intended for use in developing disaggregation methods, which can predict, from only the whole-home signal, which devices are being used. The REDD data set contains two main types of home electricity data: high-frequency current/voltage waveform data of the two power mains (as well as the voltage signal for a single phase), and lower-frequency power data including the mains and individual, labeled circuits in the house. The main directory consists of several house directories, each of which contain all the power readings for a single house. Each house subdirectory consists of a labels and channels files. The labels file contains channel numbers and a text label indicating the general category of device on this channel. Each channel.i.dat file has two columns containing UTC timestamps (as integers) and power readings (recording the apparent power of the circuit) for the channel. Experiments reported here use the House 2 data from REDD. It has 11 channels where each channel corresponds to the following appliance:

1. mains\_1
2. mains\_2
3. kitchen\_1
4. lighting

<sup>1</sup> $n=2$




Figure 2: Product of HMMs,  $P = R \times S$

5. stove
6. microwave
7. washer\_dryer
8. kitchen\_2
9. refrigerator
10. dishwasher
11. disposal

The dataset has 318759 records and 2 columns. We randomly sample 300 records for our initial experiment. Time series data from two appliances are represented as product of  $k$  HMMs.

- **Time Series** : The time series data of the microwave, dryer, kitchen\_2 and refrigerator are plotted below in Figures 3, 4, 5, 6.
- **Code** The implementation of the product of experts model is obtained from Iain Murray's website<sup>2</sup>. It implements the technique described in Geoff Hinton's paper [Hinton, 2000].
- **Additional details** Some additional details regarding experiments:
  1. The product of HMMs model (PoHMM) minimizes "contrastive divergence" as described in the paper [Hinton, 2000].
  2. The number of experts,  $k$  used here is 15. This is set somewhat arbitrarily and needs to be experimented on.
  3. Learning rate is  $\epsilon = \frac{1}{300}$ .

#### 4.4 Experimental Setup for REDD house 2

Experiments are performed on the REDD which contains 9 appliances each containing 318759 rows of energy consumption data. Experiments are done into 4 phases, in the first phase the number of data samples are varied corresponding to which the values of KL Divergence and convergence time are noted. In the second phase, the number of experts are varied keeping the best value of the sample from the first phase fixed. In the third phase, number of iterations are varied keeping the best values from above first two phases fixed. In the fourth part, the no. of appliances to be aggregated are varied.

Samples	$KLDiv$	$T(sec)$	Iterations
300	2.4864	$186.212 \pm 9.087$	18600
500	0.6761	$106.564 \pm 10.046$	10200
1000	1.1088	$158.521 \pm 1.97$	11200
1500	3.8829	$92.896 \pm 8.075$	5300
2000	1.8686	$130.98 \pm 1.932$	6900
2500	0.4733	$215.563 \pm 2.471$	9900
3000	2.8204	$258.213 \pm 1.918$	11000
3500	1.2332	$204.661 \pm 1.713$	7900
4000	0.8959	$292.666 \pm 0.619$	10400
4500	1.1118	$222.558 \pm 1.967$	7200
8000	6.392	$381.635 \pm 2.952$	8100
10000	8.276	$887.932 \pm 13.824$	10500
15000	0.7201	$1368.514 \pm 13.605$	9400

Table 4: Effect of varying samples on KL div and time

Experts	$KLDiv$	$T(sec)$	Iterations
5	0.774	$72.968 \pm 1.177$	5200
10	1.424	$117.482 \pm 1.966$	6700
15	0.473	$210.249 \pm 1.258$	9900
20	1.56	$217.739 \pm 10.452$	9000
25	7.469	$347.019 \pm 8.23$	12100
30	2.4968	$413.802 \pm 7.304$	12900
35	1.5012	$348.906 \pm 14.651$	11300

Table 5: Effect of varying experts on KL div and time

<sup>2</sup><http://homepages.inf.ed.ac.uk/imurray2/code/>




Figure 3: Microwave




Figure 4: washer\_dryer

## 4.5 Results

The evaluation of how well the learning has taken place is done by using a Kullback-Leibler divergence. KL divergence of  $P$  from  $Q$ ,  $D_{KL}(P||Q)$  is the measure of information lost when  $Q$  is used to approximate  $P$ . Here,  $P$  is the real data and  $Q$  is a fantasy data. The two probability distributions in the REDD example refer to the expert probabilities in real and fantasy data. The learned parameters from the training are fitted to the fantasy data to measure the information lost when fantasy data is used to approximate real data.

Threshold	$KLDiv$	$T(sec)$	$Iterations$
.1	0.473	$210.6 \pm 1.493$	9900
.05	0.443	$240.607 \pm 2.436$	10900
.01	0.454	$431.536 \pm 14.509$	18000
.005	0.509	$1167.243 \pm 43.412$	49800

Table 6: Effect of varying min threshold on KL div and time

Appliances	$KLDiv$	$T(sec)$	$Iterations$
3	5.559	$233.664 \pm 0.579$	10700
4	0.188	$465.634 \pm 5.275$	19900
5	.432	$338.416 \pm 3.988$	13400
6	8.736	$606.062 \pm 7.534$	28100
7	5.054	$411.457 \pm 10.051$	17300
8	0.436	$260.544 \pm 27.862$	10700
9	0.15	$474.579 \pm 14.619$	20600

Table 7: Effect of varying appliances on KL div and time

## 5 Proof of concept on Faculty housing data

### 5.1 Aim

To represent streams of energy consumption data from all the floors of faculty housing as a product of  $k$  HMMs.

### 5.2 Method

- **Data** This data represents the energy consumed by the IIIT Delhi faculty housing building. As a part of research, a team from IIIT Delhi has installed various temperature, light and motion sensors to perform real world studies and to analyse user preferences for energy conservation. For our analysis, we selected one month's historical data ranging from 01-01-2014, 00:01 hours to 31-01-2014, 23:59 hours. The two smart meters installed captures the data from all the floors. The first meter gives out readings from floors 0 to 5 and the second meter gives out readings from floors 6 to 11. The dataset includes timestamp and power consumed in watts and 84133 records. Time series data from two streams are modelled as a product of  $k$  HMMs. We also have the total power consumed by the faculty housing building which would serve as the ground truth to compare product of  $k$  HMMs with. The data is obtained from the website whose screenshot is shown in Figure 7
- **Code** It implements the technique described in Geoff Hinton's paper [Hinton, 2000].
- **Time Series** : The time series data of the energy consumption of floor 0 to 5, floor 6 to 11 and total power are plotted below in Figures 8, 9, 10.
- **Additional Details**




Figure 5: Kitchen\_2




Figure 6: refrigerator

### 5.3 Experimental Setup

Each of the data stream is modelled as a HMM individually. There are three streams of data, the first stream D1 corresponds to the data from 0-5th floor, D2 corresponds to 6-11th floor and D3 represents the total power from the faculty housing which is represented by a fixed test set, T. Firstly, the stream D1 is used to train the model such that the contrastive divergence is minimized. The parameters (mixing component of each unigauss, means of gaussian bits, log precisions of axis-aligned gaussian bits) that are learnt during the training are provided to the test set T in order to obtain the conditional probability of the gaussians given the data D1 represented as pgauss1. Similarly, the second stream of data, D2 collected from floor 6-11, is used to learn the parameters of the model during the training phase which are then again provided to the test set T to obtain conditional probability of the gaussians given the data D2 as pgauss2. Finally the data D3 is used to learn the model and parameters which are then applied to the test set T to obtain the gaussian probability as pgauss3. Now, as we know that the total power consumption of the building should be approximately equal to the product of HMMs, which is the product of pgauss1 and pgauss2. If we can show that the value of pgauss3 is as close as possible to the product of pgauss1 and pgauss2.

The experiments performed in table 8, shows the effect of varying samples on KL Divergence, convergence time and iterations keeping minimum threshold constant at 7.

The other experiment performed in table 9 shows that effect of varying experts on KL Divergence and convergence time.

### 5.4 Results

Table 8 shows that the error was minimum when the sample size was 300. With respect to the number of experts, the error was minimum when there were 5 experts as shown in table 9.

Samples	$KLDiv$	$T(sec)$	Iterations
100	2.6219e-05	257	45100
300	1.9753e-05	222	43200
500	5.5493e-05	260	44800
700	3.2847e-05	249	44000
900	3.9486e-04	221	42600
1100	4.9274e-04	317	44700
1300	3.0425e-04	276	43100
1500	3.1128e-04	303	44400
2000	1.9192e-04	306	44400
2500	1.7122e-04	370	44100
3000	1.4686e-04	331	43300
3500	1.2663e-04	370	43200
4000	1.0793e-04	403	43200

Table 8: Effect of varying samples on KL div

## 6 Conclusion & Future Work

The conclusion goes here. this is more of the conclusion

## References

- [Albert and Rajagopal, 2013] A. Albert and R. Rajagopal. Smart meter driven segmentation: What your consumption says about you. *Power Systems, IEEE Transactions on*, 28(4):4019–4030, Nov 2013.
- [Bakirtzis *et al.*, 1996] A.G. Bakirtzis, V. Petridis, S.J. Kiartzis, and M.C. Alexiadis. A neural network short term load forecasting model for the greek power system. *Power Systems, IEEE Transactions on*, 11(2):858–863, May 1996.




Figure 7: Screen shot of the webpage




Figure 8: Stream 1: Power consumption of floors 0-5

Experts	$KLDiv(e - 05)$	$T(sec)$
3	1.9780	229
4	3.5897	217
5	1.9753	228
6	4.3488	238
7	4.9111	245
8	5.6564	241
9	5.4290	258
10	5.5163	267
12	4.4504	262
14	6.9006	296
16	6.8666	300
18	6.2872	313
20	5.3842	267
25	5.8970	326
30	5.9962	327
35	5.2716	346
40	5.0955	320

Table 9: Effect of varying experts on KL div and time

- [Bassi and Olivares, 2006] Danilo Bassi and Oscar Olivares. Medium term electric load forecasting using tln neural networks. *International Journal of Computers, Communications & Control*, 1(2):23–32, 2006.
- [Brown and Hinton, 2001] Andrew Brown and Geoffrey Hinton. Proceedings of artificial intelligence and statistics 2001. In *Products of Hidden Markov Models*, number GCNU TR 2000-008, 2001.
- [Brown, 2001] Andrew Dennis Brown. Product model for sequences, 2001.
- [Chen *et al.*, 2001] Hong Chen, C.A. Canizares, and A. Singh. Ann-based short-term load forecasting in electricity markets. In *Power Engineering Society Winter Meeting, 2001. IEEE*, volume 2, pages 411–415 vol.2, 2001.
- [Chicco *et al.*, 2006] G. Chicco, Roberto Napoli, and Federico Piglion. Comparisons among clustering techniques for electricity customer classification. *Power Systems, IEEE Transactions on*, 21(2):933–940, May 2006.
- [Chow and Leung, 1996] T.W.S. Chow and C.T. Leung. Nonlinear autoregressive integrated neural network model






Figure 9: Stream 2: Power consumption of floors 6-11




Figure 10: Total Power of the building

for short-term load forecasting. *Generation, Transmission and Distribution, IEE Proceedings-*, 143(5):500–506, Sep 1996.

[Daneshi *et al.*, 2008] H. Daneshi, M. Shahidehpour, and A.L. Choobbari. Long-term load forecasting in electricity market. In *Electro/Information Technology, 2008. EIT 2008. IEEE International Conference on*, pages 395–400, May 2008.

[Fabre *et al.*, 2000] Eric Fabre, Stefan Haar, Albert Benveniste, and Albert Benveniste. Hidden markov models for distributed and concurrent systems, 2000.

[Falvo *et al.*, 2006] M.C. Falvo, R. Lamedica, S. Pierazzo, and A. Prudenzi. A knowledge based system for medium term load forecasting. In *Transmission and Distribution Conference and Exhibition, 2005/2006 IEEE PES*, pages 1291–1295, May 2006.

[Fan and McDonald, 1994] J.Y. Fan and J.D. McDonald. A real-time implementation of short-term load forecasting for distribution power systems. *Power Systems, IEEE Transactions on*, 9(2):988–994, May 1994.

[Ghahramani and Jordan, 1997] Zoubin Ghahramani and Michael I. Jordan. Factorial hidden markov models. *Mach. Learn.*, 29(2-3):245–273, November 1997.

[Heinzelman *et al.*, 2000] Wendi Rabiner Heinzelman, Anantha Ch, and Hari Balakrishnan. Energy-efficient communication protocol for wireless microsensor networks. pages 3005–3014, 2000.

[Hinton, 2000] Geoffrey E. Hinton. Training products of experts by minimizing contrastive divergence. Technical Report GCNU TR 2000-004, Gatsby Computational Neuroscience Unit, University College London, 2000.

[Kim *et al.*, 2011] Hyungsul Kim, Manish Marwah, Martin F. Arlitt, Geoff Lyon, and Jiawei Han. Unsupervised disaggregation of low frequency power measurements. In *SIAM*, pages 747–758. SIAM / Omnipress, 2011.

[Kolter and Ferreira, 2011] J. Zico Kolter and Joseph Ferreira. A large-scale study on predicting and contextualizing building energy usage. In *AAAI*. AAAI Press, 2011.

[Kolter and Jaakkola, 2012] J. Zico Kolter and Tommi Jaakkola. Approximate inference in additive factorial hmms with application to energy disaggregation. In *AISTATS*, volume 22 of *JMLR Proceedings*, pages 1472–1482. JMLR.org, 2012.

[Martinez Alvarez *et al.*, 2011] F. Martinez Alvarez, A. Troncoso, J.C. Riquelme, and J.S. Aguilar Ruiz. Energy time series forecasting based on pattern sequence similarity. *Knowledge and Data Engineering, IEEE Transactions on*, 23(8):1230–1243, Aug 2011.

[McKerracher and Torriti, 2013] Colin McKerracher and Jacopo Torriti. Energy consumption feedback in perspective: integrating australian data to meta-analyses on in-home displays. *Energy Efficiency*, 6(2):387–405, 2013.

[Mirowski *et al.*, 2014] Piotr Mirowski, Sining Chen, Tin Kam Ho, and Chun-Nam Yu. Demand forecasting in smart grids. *Bell Labs Technical Journal*, 18(4):135–158, 2014.

[Samuel *et al.*, 2014] Isaac A Samuel, Ayokunle A Awelewa, et al. Medium-term load forecasting of covenant university using the regression analysis methods. *Journal of Energy Technologies and Policy*, 4(4):10–16, 2014.

[Stern, 2011] David I. Stern. The role of energy in economic growth. *Annals of the New York Academy of Sciences*, 1219(1):26–51, 2011.



- [Taylor and Hinton, 2009] Graham W. Taylor and Geoffrey E. Hinton. Products of hidden markov models: It takes n  $\geq$  1 to tango. In *Proceedings of the Twenty-Fifth Conference on Uncertainty in Artificial Intelligence, UAI '09*, pages 522–529, Arlington, Virginia, United States, 2009. AUAI Press.
- [Wijaya *et al.*, 2013] Tri Kurniawan Wijaya, Julien Eberle, and Karl Aberer. Symbolic representation of smart meter data. In *Proceedings of the Joint EDBT/ICDT 2013 Workshops*, EDBT '13, pages 242–248, New York, NY, USA, 2013. ACM.
- [Wijaya *et al.*, 2014] Tri Kurniawan Wijaya, Tanuja Ganu, Dipanjan Chakraborty, Karl Aberer, and Deva P Seetharam. Consumer segmentation and knowledge extraction from smart meter and survey data. In *SIAM International Conference on Data Mining (SDM14)*, number EPFL-CONF-196276, 2014.
- [Zhang and Ye, 2011] Zhiheng Zhang and Shijie Ye. Long term load forecasting and recommendations for china based on support vector regression. In *Information Management, Innovation Management and Industrial Engineering (ICIII), 2011 International Conference on*, volume 3, pages 597–602, Nov 2011.