

BODHI: A System for Crowdsourced OCR Text Correction

Student Name: Megha Gupta

IIIT-D-MTech-CS-DE-13-MT11024

Indraprastha Institute of Information Technology, New Delhi
meghag@iiitd.ac.in

Advisor(s): Dr. Haimonti Dutta and Mr. Manoj Pooleery
The Center for Computational Learning Systems (CCLS),
Columbia University, New York.
{haimonti,manoj}@ccls.columbia.edu

Submitted in partial fulfillment of the requirements
for the degree of MTech in Computer Science
with Specialization in Data Engineering

July 22, 2013

1 Abstract

Newspapers are the first draft of history – they are a rich source of information for historians, researchers, and scholars. With the advent of *digitized* newspapers, the accessibility to old historic papers has increased. The usability of archives storing these newspapers depends on the imaging technology, Optical Character Recognition (OCR) devices, zoning and segmentation, metadata extraction, searchability and web delivery systems developed to make them accessible. It is a well known fact that the OCR technology is far from perfect and often, the text generated is garbled affecting the efficiency of search and retrieval. Since the human mind excels in visual cognition and language processing tasks and machines are not able to completely recreate these skills, manual labor is involved in correcting garbled OCR.

Our goal in this project is to build a system that will allow users of a digitized newspaper archive to read and correct OCR text. This report is organized as follows: Section 2 and 3 describe the functions and workflow of the system, Section 4 and 5 elaborates on its design model and System specifications; finally, implementation details with real screenshots and the challenges faced during development are mentioned in Section 6.

Contents

1	Abstract	2
2	Introduction	4
2.1	Problem Description	4
2.2	Scope & Objectives	4
2.3	Software Functions	4
3	Software Project Plan	5
3.1	Workflow of the System	5
4	Design Models	7
4.1	Database Design	7
4.2	Architectural Design	8
4.2.1	Architectural Description	9
5	System Study	10
5.1	Hardware Specification	10
5.2	Software Specification	10
5.3	Setting up the code base	11
6	Implementation	13
6.1	Functions Implemented	13
6.2	User Interface	14
6.3	Remote Debugging Tomcat using Eclipse	17
6.4	Challenges	18
7	Conclusion	19
8	Acknowledgement	19

2 Introduction

2.1 Problem Description

The goal of our project is to develop a system to edit garbled OCR text from old historic newspapers. The archive under consideration is the holdings of California Digital Newspaper Collection (CDNC) [1]. The OCR software used on digital images of newspapers in this archive has generated garbled text due to poor quality of input paper¹, varying size and style of font and different column layouts. This adversely affects the retrieval effectiveness, hence the techniques for cleaning the OCR [2] need to be improvised. Often such techniques involve laborious and time consuming manual processing of data. By correcting the garbled text, we hope to be able to improve the retrieved results thus making the archive more useable.

2.2 Scope & Objectives

The objective of this project is to build a a web interface equipped with a text correction tool for enabling the users to correct the OCR errors as they come across them. The software is expected to be usable, intuitive, simple and functions well consistently. The basic functions are:

1. Searching and displaying images of issues
2. Search by issue name and date range
3. User management functions
4. Text correction functionality

2.3 Software Functions

The project functions can be categorized into the following:

- User management functions
These encompass all the activities to manage the users for the application. These include:

¹Newspapers date back to 1846

1. Registration module: Ability to enter and validate user demographic information
 2. User Security module: Ability to authenticate and authorize the user
 3. User tracking module: Ability to track user activity, prepare statistics (how long users spent on each module, how many articles they corrected, participation in discussions etc)
- Article related functions
 1. Articles search: This provides the capability to retrieve an article based on a set of criteria like article topic, date range etc.
 2. OCR correction: This interface provide users with the capability to correct the OCR displayed along side a given article image. A basic level of authentication (based on a CAPTCHA model) will be required before the user can correct the text.
 3. Tag specification: This interface will allow the user to specify tags corresponding to a given article. This functionality will also require a basic level of authentication.
 - Application administration functions

3 Software Project Plan

This project was divided into three phases:

Phase 1 included laying out UI components, database connectivity, searching and displaying user requested issues with their images.

Phase 2 included basic functionality text correction tool, user management functions.

Phase 3 included advance searching and functionality of tool.

3.1 Workflow of the System

- In Phase 1, the first task was to design and build the layout of the interface. The user interface components in Vaadin can roughly be divided in two groups: components that the user can interact with and layout

components for placing the other components to specific places in the UI. We started by creating a content layout for the UI, and then added the other layout components hierarchically, and finally the interaction components as the leaves of the component tree.

The Root Layout in this project is the Vertical Layout which encompasses two Horizontal Layouts, called `mainLayout` for the search options and the other called as `bottomLayout`. The bottomLayout is again divided into three panels called tool panel, image panel and function panel at left, center and right side of the screen respectively. The tool panel consists of tables, text correction tool, etc. The image panel consists of newspaper's images. And the function panel contains the user management controls.

- Then we had to incorporate the basic search functionality on the basis of issue names given in the combobox, using date range from the date picker. Once a issue is selected from the dropdown list and the search button is clicked, we get all the headlines of the selected issue in a table in the tool panel. And the images of the corresponding issue is displayed in the image panel. The headlines are in the form of links which when clicked shows the entire article in the table itself.
- In Phase 2, the actual implementation of collaborative text correction tool takes place. Once we get the entire OCR text in the rows of a table, it also gets highlighted in the image at the image panel. Once `Edit Text` button is clicked, the OCR text becomes editable and a save button is displayed at the bottom of the left hand pane. A user can correct the text in the left panel to match it with the text in the image.
- Once the `Save` button is clicked, a check is performed against the database to determine which lines of text have been changed. Only those text lines that have been changed are saved as a new version.
- In Phase 3, we incorporate user management functions, advance search options like phrase searching, enhance UI features and other functionalities of the tool.

4 Design Models

4.1 Database Design

The database design of our project is as given below.

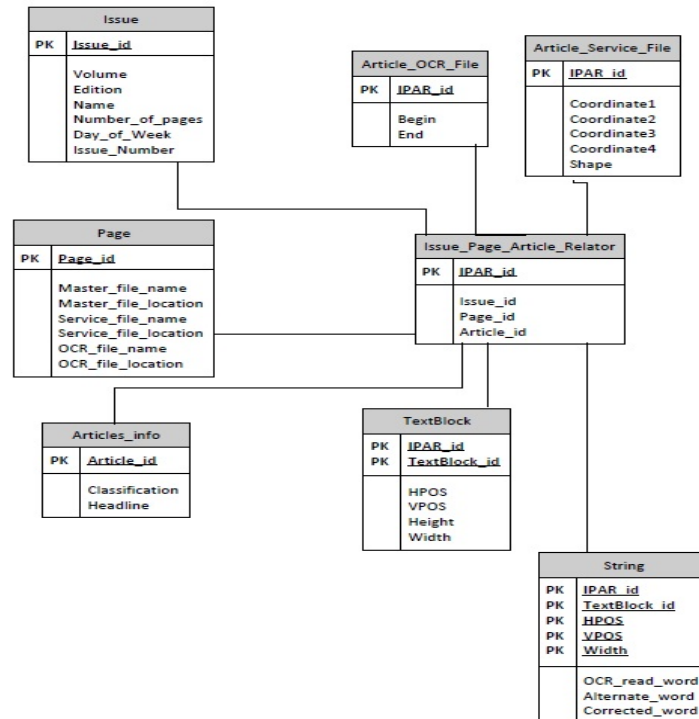


Figure 1: Database Design

Following are the steps to access the Database Server (Postgres) using a command line client:

1. To connect to the database on Tesla.ldeo.columbia.edu, a tunnel needs to be established. To do this, type
ssh -Nn -L 5555:localhost:5432 username@tesla.ldeo.columbia.edu &
where username is the corresponding user's name.
2. Open psql command line client and type in the following details: Server
[localhost] : localhost

Database [postgres]: nypltest
Port [5432]: 5555
Username [postgres]: nypl_user
password for nypl_user: c9MamuaT

3. Commands like `\l` and `\d` gives the list of Databases and Relations in PostgreSQL.

We have used views to access the tables present in the database. A View is a virtual table containing fields from one or more real tables. They make information retrieval fast and secure.

Following Views are used by our front end:

- `vw_issu_page_files`
To get the page file locations (and related data) related to particular issue using the `'date_of_issue'` column.
- `vw_issue_articles`
To access all the articles related to a particular issue using the `'date_of_issue'` column.
- `vw_image_info`
To access the coordinates, shapes and related data for each article using the `'ipar_id'` column.
- `vw_article_text`
To access entire text related to a particular article using the `'ipar_id'` column

4.2 Architectural Design

The architectural style that suits this project is the Model-View-Controller (MVC) architecture. It is the software architecture pattern that separates the representation of information from user's interaction with it. It is used by applications that need the ability to maintain multiple views of the same data. The interaction between the three components is defined as:

- A controller sends commands to its associated view to change the view's presentation of the model. It can also send commands to the model to update the model's state. It processes every request, prepares other parts of the system like model and view.

- A model handles data processing and database works part. It processes events sent by controller. After processing these events, it sends processed data to controller (thus, controller may reprocess it) or directly to view side allowing them to produce updated output.
- A view prepares interface to show to the user. It requests from the model the information that it needs to generate an output representation to the user.




Figure 2: MVC Architecture

4.2.1 Architectural Description

In our project, we have MainWebApp.java as the controller which receives calls from the user and calls the relevant screen (model or view) based on the input.

Then we have IssueDisplay.java which acts as a view. This class is the primary display page for an issue, one page at a time. It takes in the following parameters - list of articles in an issue and the page image to be displayed first. It then display the list of articles as hyperlinks in the left pane, along with the issue title and display the image in the right pane.

Then we have DatabaseHelper.java which acts as a model containing any database objects that needs representation.

5 System Study

5.1 Hardware Specification

- 32-bit Operating System
- 4GB RAM
- Intel(R) Core(TM) i3 CPU @ 2.13 GHz

5.2 Software Specification

- Eclipse Juno 4.2 IDE
It is an Integrated Development Environment (IDE) comprising of a base workspace and an extensible plug-in system for customizing environment. It is used to develop applications primarily in Java and other languages by means of other plug-ins.
We used the Vaadin plug-in to create our project in Eclipse Juno.
- Vaadin 6 (Java Framework) plugin for Eclipse
Vaadin is an AJAX web application development framework that enables developers to build high-quality user interfaces with Java, both on the server- and client-side. It provides a set of libraries of ready-to-use user interface components and a clean framework for creating our own components. The focus is on ease-of-use, re-usability, extensibility, and meeting the requirements of large enterprise applications.
It is an open source web application for rich Internet applications. It has server side architecture which means that the majority of the logic runs on the server. Ajax technology is used at the browser side to ensure a rich and interactive user experience. On the client-side Vaadin is built on top of Google Web Toolkit (GWT).
- Apache Tomcat 7.0.35 It is an open source web server and servlet container developed by the Apache Software Foundation (ASF). It implements the Java Servlet and the JavaServer Pages (JSP) specifications

from Sun Microsystems, and provides a pure Java HTTP web server environment for Java code to run in.

It includes tools for configuration and management, but can also be configured by editing XML configuration files.

- Apache Ant 1.8.4 It is a software tool for automating software build processes. It is similar to `Make` but is implemented using the Java language, requires the Java platform, and is best suited to building Java projects.

The most immediately noticeable difference between Ant and Make is that Ant uses XML to describe the build process and its dependencies, whereas Make uses Makefile format.

- PostgreSQL It is an open source object-relational database management system (ORDBMS) available for many platforms including Linux, Microsoft Windows and Mac OS X.

It implements the majority of the SQL:2008 standard, is ACID-compliant, is fully transactional (including all DDL statements), has extensible data types, operators, index methods, functions, aggregates, procedural languages, and has a large number of extensions written by third parties.

- Tortoise SVN 1.7, A Subversion client for Windows It is a free software that helps programmers manage different versions of the source code for their programs. Tortoise SVN is a Subversion client, implemented as a Microsoft Windows shell extension. It is released under the GNU General Public License.
- Web Browser, Google Chrome or Mozilla Firefox

5.3 Setting up the code base

Following are the steps to get the application working:

1. Checkout the Bodhi Webapps code located here -
<https://power.ldeo.columbia.edu/svn/Proj/NYPL/Bodhi/Code/WebApps>
using Twiki username and password.
2. Verify that it has the following directory structure:

- build (directory where the classes of some components will be located, for now, this is empty)
 - config (This folder should have two files - log4j.properties and BodhiWebApplication.properties)
Edit log4j.properties, change the value for `log4j.appender.A1.File` to a specific directory on the machine.
No changes are required in BodhiWebApplication.properties
 - doc (This is the directory for any documents, including javadoc)
 - resources (If the application is going to use any external resources, they will be stored here)
 - src (Root folder for all the source files)
 - edu.columbia.ccls.dh.bodhi.controller (Folder for all the controllers, these will have the logic to direct the calls to specific components. It will have the code to instantiate the models and views corresponding to the action chosen by the user)
 - edu.columbia.ccls.dh.bodhi.model (Folder for all the model classes, typically any database objects that need representation)
 - edu.columbia.ccls.dh.bodhi.view (Folder for constructing the HTML views. Information needed for constructing will be passed to it by the constructor, typically either an instance of the model or specific values from the model)
 - edu.columbia.ccls.dh.bodhi.utils (Folder for all the utility classes and functions. Components in this folder will be used across the board, like database connections, special string operations, validations etc.)
 - WebContent (Root folder for all the Vaadin files as well as the web application classes)
3. Set up the project in Eclipse, just by choosing create new project from existing source and pointing it to the root folder where the entire folder is checked out.
 4. Once the project is created, make sure that the java path is set correctly (it might be different on different machines)

5. Once the project is ready, open the deploy.xml, locate the entry , change the value of the attribute `location` , point it to the Tomcat installation directory on the machine.
6. Make sure Apache Ant is installed on the machine (if not, download it from apache.org and make sure that the executable location is included in the PATH variable)
7. Verify that there are no errors, then open up a command prompt
8. To connect to the database on Tesla.ldeo.columbia.edu, a tunnel needs to be established. To do this, type
ssh -Nn -L 5555:localhost:5432 username@tesla.ldeo.columbia.edu &
Replace the username and password with the corresponding user's name and password.
9. Open another command prompt, change the directory to the root folder of the project and type ant war. This will compile the project, create a war file and deploy it to the Tomcat's `webapps` directory.
10. Start Tomcat (if it is not already running), go to bin and run startup.bat to start the server and shutdown.bat to shutdown the server.
11. Open up a web browser, go to http://localhost:8080/bodhi/MainWebApp, to see the page, Bodhi Issue Display created by the application.
12. To verify that the application is up and running correctly, open up the log file that was specified in step number 2

6 Implementation

6.1 Functions Implemented

1. List of Issues : Populated the combobox with the retrieved list of issues from the database. See figure 1
2. Article Search : Retrieved an article based on a set of criteria like article topic, date range. See figure 3
3. Headlines : Retrieved articles are the headlines in the form of links which when clicked returns the OCR text. See figure 3

4. Images: Retrieved images of the selected articles in the middle pane.
5. OCR Text : Displayed the raw OCR article text of the user selected headline in the table.
6. Edit Raw OCR text : Raw OCR text can be edited after clicking the `Edit` button. Once the button is pressed, each row of the table becomes editable.
7. Save Edited Text : Text edited by the patron gets stored in the database once the `Save` button is pressed.

6.2 User Interface

The User Interface is designed using Vaadin. Some of our project screenshots are as shown below.

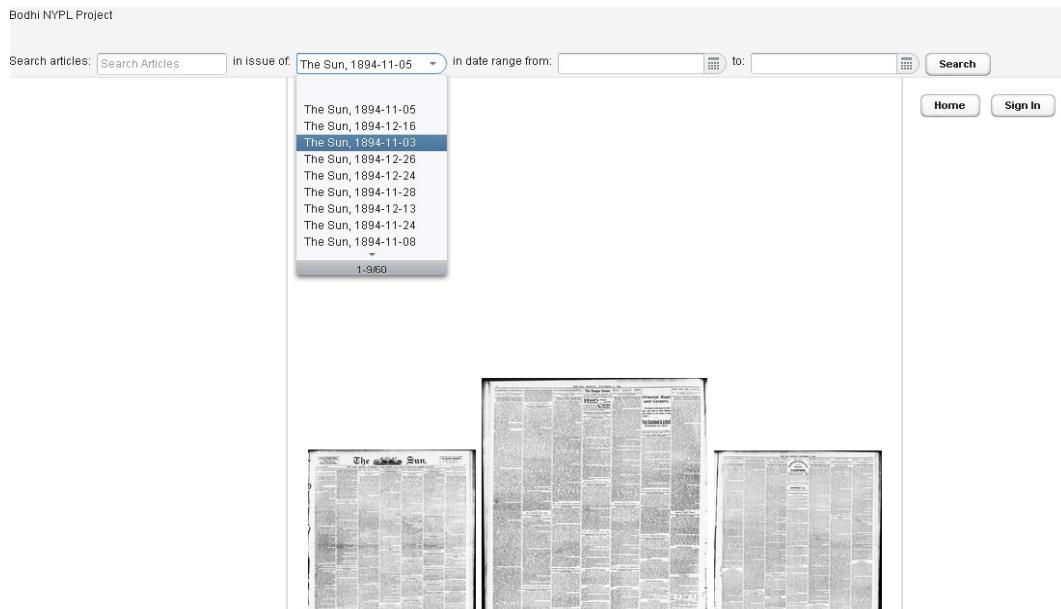


Figure 3: List of issues




Figure 4: Datepicker



Figure 5: List of headlines in a issue




Figure 6: OCR Text of the headline 'Amusements'

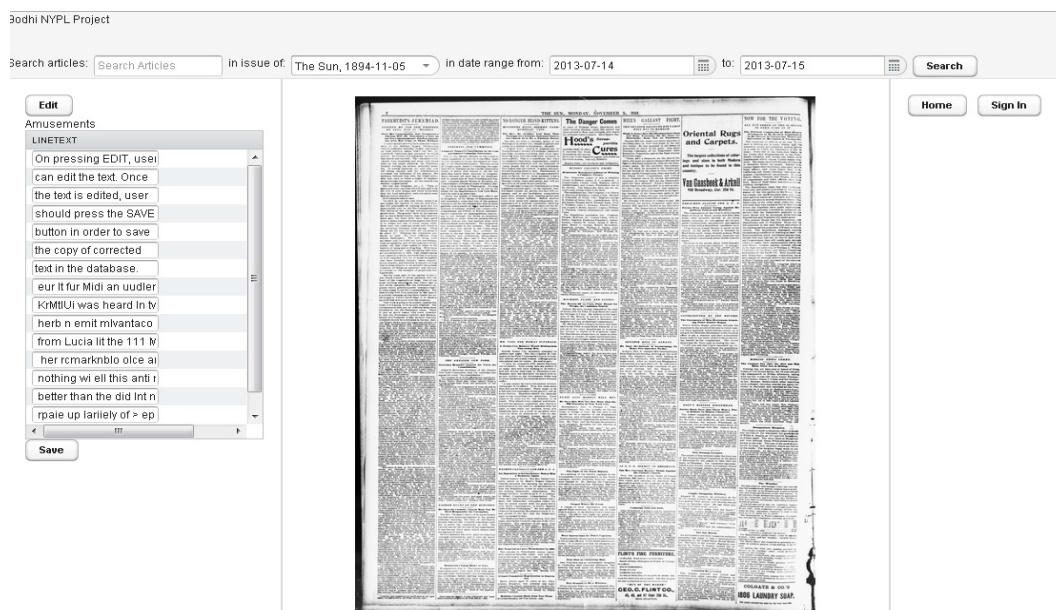


Figure 7: Editing the Raw Text

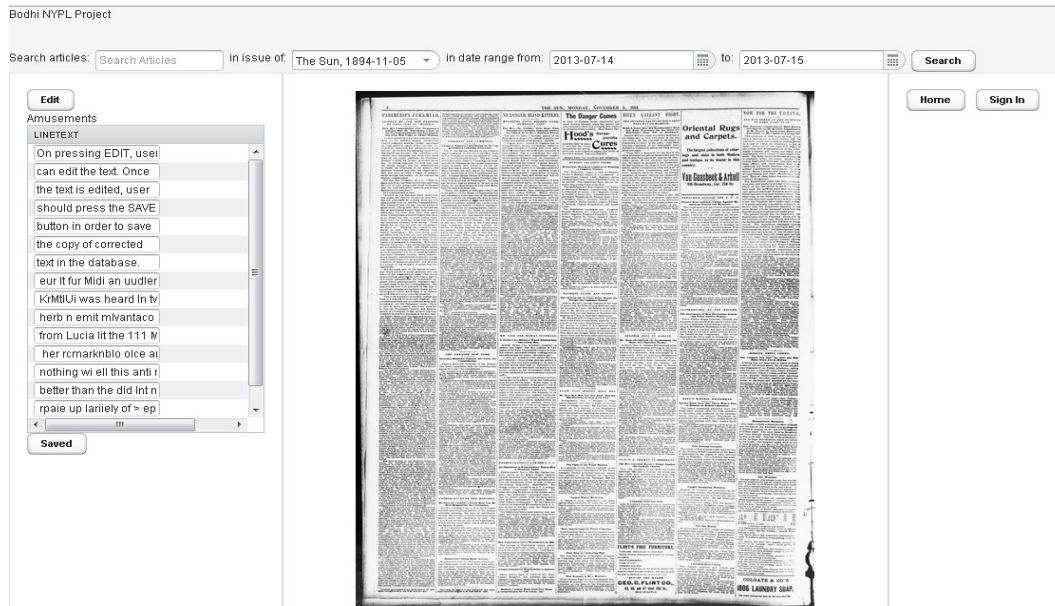


Figure 8: Edited Text Saved

6.3 Remote Debugging Tomcat using Eclipse

To debug our project, we could not simply use the Eclipse Integrated debugger as to debug an executable file is required. So in order to debug a Tomcat application through Eclipse, following steps need to be taken:

1. Open catalina.bat file which can be found in apache-tomcat-7.0.35/bin/catalina.bat
2. Insert two statements after "set CLASSPATH="


```
CATALINA_OPTS="-Xdebug-Xrunjdwp:transport=dt_socket,address=8000,server=y,suspend=y"
      JPDA_OPTS="-agentlib:jdwp=transport=dt_socket,address=8000,server=y,suspend=n"
```
3. Save and close the file.
4. Run Tomcat from command prompt using command:


```
catalina.bat jpda start
```

 instead of using the normal startup.bat
5. Then in Eclipse create a debug configuration like below:
 - Give the project name.

- Give the connection type as Standard(Socket Attach)
- Put the host as localhost
- Specify port as 8000
- Give any name for configuration.

6. After setting the debug configuration, start debugger in Eclipse.

6.4 Challenges

There were several challenges faced during the project, few of them are listed below.

1. Firewall at server end: Since the database server was at an external location secured with firewall, there were problems in accessing the server at first.
2. Firewall at client end: After the access was granted the problem arose at my end due to Cyberoam layer. I could access the server from outside but not institution due to restricted remote access in the presence of firewall. Finally I was allotted a port using which I could successfully create a secure tunnel.
3. Inaccessible Images: Images that were to be displayed in the web page were not available publicly as they were present only on the server. So hundreds of images of size 3.5GB were first securely copied (PSCP) to the client and then they were pushed to the cloud.
4. Incompatible Images: All the images that were available were of the format .jp2 which is JPEG 2000. It uses a wavelet compression algorithm instead of the common DCT compression method that is used with standard JPEG images.
This format can be used to store either lossy or lossless compressed JPEG images and provides the user with a better quality and smaller file size than the traditional JPG image files.
But these images are not supported by the browser, hence they can not be opened in the browser.
5. Performance Issues : Due to the remote access to the server through tunneling and by using complex subquery to retrieve the OCR text,

the server got slow. The time clocked for the server to response for the query was 16 minutes when using 3G network. These long delays and other problems made it difficult to work on real data.

So, the functionality to edit the Raw text was built by setting OCR text in the table a priori. And saving the edited data into the database.

7 Conclusion

The OCR [2] Text Correction System has not been fully implemented yet. There are several challenges as described above that require attention before proceeding further in terms of implementation. This system when functional will contain a repository of articles which would be clean and usable by many. The process of editing text will also help in building and growing of a community.

Once the system would be functional, it will certainly act as a rich source of information for historians, researchers, scholars and users in general.

8 Acknowledgement

This implementation project was a part of my M.Tech Scholarly work. I would like to express my gratitude to both my Advisors Dr. Haimonti Dutta and Manoj Pooleery for their constant support and guidance without which I would have never been able to complete it.

References

- [1] California Digital Newspaper Collection. <http://cdnc.ucr.edu/cdnc>, June 2009.
- [2] Line Eikvil. Optical character recognition. 1993.
- [3] Marko Grnroos. *Book of Vaadin: Vaadin 7 Preview edition*. Vaadin Ltd, 2013.
- [4] Allen Condit Kazem Taghva, Julie Borsack. *Evaluation of Model Based Retrieval Effectiveness with OCR Text*, volume 14. New York, US, 1996.

- [5] Allen Condit Kazem Taghva, Julie Borsack and Srinivas Erva. *Effects of Noisy Data on Text Retrieval*, volume 45. Las Vegas, US, 1993.